

## Disk interfaces

- Relatively few
  - SCSI (pronounced "scuzzy")
    - Common, widely supported
  - IDE a.k.a. ATA or PATA, and SATA
    - Inexpensive, simple
  - Fibre Channel
    - High bandwidth, lots of simultaneous devices
    - Supports up to 16Gbit
  - Universal Serial Bus (USB)
    - Typically used for slow devices (e.g., CD-ROMs, portable, removable drives)

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## CSE 265: System and Network Administration

- Disks
- Partitions
- Volumes
- Filesystems
- Files




Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## IDE a.k.a. ATA

- Integrated Drive Electronics / AT Attachment
  - Very short cable lengths (18in!)
- ATA-2 added DMA and LBA (get beyond BIOS 504MB limit)
- ATA-3 added power management, self-monitoring (16MB/s)
- Ultra-ATA added Ultra DMA/33, /66, and /133 modes (33-133MB/s)
- Hard disks with this interface were last produced in 2013
- ATAPI interface allows non-ATA devices to connect
  - E.g., CD-ROMs




Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## SCSI: Small Computer Systems Interface

- Many versions
  - SCSI-1 (1986) 8-bits, 5MB/s
  - SCSI-2 (1990) added command queuing, DMA, more
  - Fast SCSI-2 8-bits, 10MB/s
  - Fast/wide SCSI-2 16-bits, 20MB/s
  - Ultra SCSI 8 bits, 20MB/s
  - Wide Ultra SCSI 16bits, 40MB/s
  - Wide Ultra2 SCSI 16bits, 80MB/s
  - Wide Ultra3 SCSI 16bits, 160MB/s
  - Ultra-320, Ultra-640 SCSI




Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Black box

- 40+2 SATA drives
- RAID
- Dual Xeon
- 8U tall
- Up to 80TB



Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison



Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## SATA

- Now standard equipment
  - Fast: 150-600MB/s (16GBit/s now available)
  - Software compatible with parallel ATA
  - One drive per controller
  - Thin cables



## Sun X4500

- 48 SATA drives
- Software RAID or ZFS
- Dual AMD
- 4U tall
- Up to 48TB



Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## SCSI vs. SATA

- SCSI traditionally beats SATA technically, but may not be worth the price premium
  - In single-user systems, SATA will provide 85%, cheaply
- For best possible performance, SCSI is often better
  - e.g., in servers and multi-user systems
  - handles multiple simultaneous reqs + more devices better
  - higher-end equipment (faster, better warranty, etc.)
- SATA technology is quite good
  - Usually better price/performance than SCSI
- Still subject to much debate

# Disk partitions

- Drives are divided into one or more partitions that are treated independently
  - Partitions make backups easier, confine damage
- Typically have at least two or three
  - root partition (one)
    - everything needed to bring system up in single-user mode (often copied onto another disk for emergencies)
  - swap partition (at least one)
    - stores virtual memory when physical memory is insufficient
  - user partition(s)
    - home directories, data files, etc.
  - boot partition - boot loader, kernel, etc.

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

# Adding a disk to Linux

## STEP-BY-STEP (w/out LVM)

- Install new hardware
  - verify that hardware is recognized by BIOS or controller
- Boot, make certain device files already exist in /dev
  - e.g., /dev/sdc
- Use **fdisk/parted** (or similar) to partition the drive
  - Verify the system type on each partition
- Use **mke2fs** (-t ext4) on each regular partition
  - To create (an ext4) filesystem
- Use **mkswap** to initialize swap partitions
- Add entries to /etc/fstab
- Mount by hand, then reboot to verify everything

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

# Logical Volumes

- Partitions are static, and sometimes you want to change them
- LVM (Linux Logical Volume Manager) lets you combine partitions and drives to present an aggregate volume as a regular block device (just like a disk or partition)
  - Use and allocate storage more efficiently
  - Move logical volumes among different physical devices
  - Grow and shrink logical volume sizes on the fly
  - Take “snapshots” of whole filesystems
  - Replace on-line drives without interrupting service
- Similar systems are available for other OSes

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

# hdparm: test/set hd params

- hdparm will do simple performance tests

```
[root@wume2 ~]# /sbin/hdparm -Tt /dev/hda
/dev/hda:
Timing cached reads:   1928 MB in  2.00 seconds = 963.26 MB/sec
Timing buffered disk reads: 122 MB in  3.03 seconds = 40.22 MB/sec

[root@wume1 ~]# /sbin/hdparm -Tt /dev/sda
/dev/sda:
Timing cached reads:   3440 MB in  2.00 seconds = 1720.77 MB/sec
Timing buffered disk reads: 162 MB in  3.03 seconds = 53.41 MB/sec

[root@night ~]# /sbin/hdparm -Tt /dev/sdd
/dev/sdd:
Timing cached reads:   10504 MB in  2.00 seconds = 5254.65 MB/sec
Timing buffered disk reads: 1196 MB in  3.00 seconds = 398.28 MB/sec

[root@morning ~]# /sbin/hdparm -Tt /dev/hda
/dev/hda:
Timing cached reads:   4092 MB in  2.00 seconds = 2047.82 MB/sec
Timing buffered disk reads: 10 MB in  3.03 seconds = 3.30 MB/sec
```

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## ext# filesystems

- For ext2/ext3/ext4, **mke2fs** is used, which creates
  - A set of inode storage cells
    - each holds info about one file
  - A set of scattered “superblocks”
    - holds global filesystem info (multiple copies for reliability)
    - size and location of inode tables, block map and usage, etc.
  - A map of the disk blocks in the filesystem (used and free)
  - The set of data blocks

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison


Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## LVM

- Sample organization:



## Mounting a filesystem

- Filesystem must be mounted before use
  - Must be made part of root filesystem
- Can be mounted on (top of) any directory
  - # **mount /dev/sda1 /usr/local**
  - # **df /usr/local**
- Use /mnt for temporary mounts
- Want to set up automatic mounting

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Filesystems

- Linux filesystems are created in partitions or volumes
  - ext2fs (2nd Extended File System) is old
  - ext3fs (3rd Extended File System) is common
    - Augments ext2fs to incorporate journaling
      - Journals contain filesystem updates
      - Journal log can reconstruct consistent filesystem
      - Journal speeds filesystem consistency checks
  - ext4fs (Fourth Extended File System) is modern
    - Speeds large directories
    - Compatible with ext2 and ext3
  - Other filesystems also supported
    - ReiserFS, IBM's JFS, SGI's XFS
  - Can read foreign filesystems (e.g., FAT, NTFS, ISO 9660)

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## fsck: check and repair filesystems

- During power failure, superblock, inodes, and data blocks may not get written to disk
- **fsck** can fix minor damage (ext3/4 systems quickly)
  - unreferenced inodes
  - inexplicably large link counts
  - unused data blocks not recorded in block maps
  - data blocks listed as free that are also used in a file
  - incorrect summary info in superblock
- More complex damage will make **fsck** ask human
  - Places unfixable files in lost+found directory
  - You should re-run **fsck** until no errors are found

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## /etc/fstab

- (Almost) every filesystem that the system knows about automatically is in /etc/fstab

```
[root@brian]# more /etc/fstab
# /etc/fstab
# Created by anaconda on Thu Jan 19 14:11:35 2012
#
# Accessible filesystems, by ref., are maintained under '/dev/disk'
# See man pages fstab(5), findfs(8), mount(8) and blkid(8) for more
#
/dev/mapper/vg_davison-lv_root / ext4 defaults 1 1
UUID=52bb6031-5fda-402e-bb9f-5c0fee93ca44 /boot ext4 defaults 1 2
/dev/mapper/vg_davison-lv_home /home ext4 defaults 1 2
/dev/mapper/vg_davison-lv_swap swap swap defaults 0 0
tmpfs /dev/shm tmpfs defaults 0 0
devpts /dev/pts devpts gid=5,mode=620 0 0
sysfs /sys sysfs defaults 0 0
proc /proc proc defaults 0 0
```

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## The Filesystem

- A filesystem incorporates:
  - A way of naming and organizing things (namespace)
  - An API for navigating and manipulating objects
  - A security model for protecting, hiding, and sharing objects
  - An implementation to tie the model to the hardware
- Linux abstract kernel interface supports many different filesystems
  - from disk, network, memory

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## [u]mounting, swap

- **mount**, **umount**, **swapon** and **fsck** all read the /etc/fstab file
- enables
  - # **mount /mnt/cdrom**
- fstab entries must be in the correct order
- at startup
  - **mount -a** executed, mounts all regular partitions
  - **swapon** enables swapping on all swap partitions

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison



## [u]mounting filesystems

---

- List of filesystems is in `/etc/fstab`
  - Such filesystems are checked (**fsck -A**) and mounted (**mount -a**) at boot
- `umount` will fail if the filesystem is busy
  - `busy` = any open files, processes with `cwd`, or copies of executing programs
  - **/sbin/fuser** will show such processes
    - `f` – file open for reading or writing
    - `c` – process `cwd` is on filesystem
    - `e` – process is executing a file
    - `r` – process `root dir` is on filesystem
    - `m` – process has mapped file or shared lib

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Pathnames

---

- The Linux filesystem is a single unified hierarchy, starting with `/` (the root directory)
- A pathname can be
  - absolute
    - `/etc/passwd`
  - relative
    - `./passwd`
    - Always starts with current working directory
- No technical limitations on file naming other than length and `/`
  - some chars are more difficult to use (need quotes or escape)

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## File tree organization

---

- Not really well organized
- Many files organized by function
  - difficult to upgrade
  - `/etc/` contains files that are never customized, and ones that are entirely local
- There is at least one place for everything
- Admins need to learn standard places, not move or use new ones

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Mounting & unmounting filesystems

---

- The filesystem is made of smaller filesystems
- Most filesystems occupy disk partitions
  - but can be anything that obeys the API
- Filesystems may be added or removed using the `mount` and `umount` commands
  - The mount point is a directory
  - Ex: **# mount /dev/hdc1 /backup**

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

# File types

- Linux defines seven types of files

- `[-]` - Regular files
  - `[d]` - Directories
  - `[c]` - Character device files
  - `[b]` - Block device files
  - `[s]` - Local domain sockets
  - `[p]` - Named pipes (FIFO)
  - `[l]` - Symbolic links

- **ls -ld** shows the filetype of a file

# Directories

- Created with **mkdir**, deleted with **rmdir** (if empty) or **rm -r**
- Contains named references (links) to other files
- Special entries “.” and “..” refer to self and parent directories respectively
- Filenames are stored within parent directory
- More than one directory entry can refer to the same file (hard links)
  - Can be created with **ln**, removed with **rm**

# Filesystem hierarchy

<http://www.pathname.com/fhs/>

`/bin` : Essential user command binaries (for use by all users)  
`/boot` : Static files of the boot loader (e.g., kernel)  
`/dev` : Device files (terminals, disks, modems, etc.)  
`/etc` : Host-specific system configuration  
`/home` : User home directories (optional)  
`/lib` : Essential shared libraries and kernel modules  
`/media` : Filesystems on removable media  
`/opt` : Add-on application software packages  
`/proc` : Kernel and process information virtual filesystem  
`/root` : Home directory for the root user (optional)  
`/sbin` : Static system binaries for repairing, booting, & recovering OS  
`/tmp` : Temporary files (that disappear at reboot)  
`/usr` : (more next slide)  
`/var` : (more next slide)

# /usr, /var

`/usr`  
`/usr/bin` : Most commands and executables  
`/usr/include` : Header files for C programs  
`/usr/lib` : Libraries and support files for standard programs  
`/usr/local` : Local software (stuff you install)  
`/usr/man` : Manual pages  
`/usr/sbin` : Less essential sysadmin commands  
`/usr/share` : Content that is common to multiple systems (RO)  
`/usr/src` : Source code for (nonlocal) software packages  
`/var`  
`/var/adm` : Various logs, system setup records  
`/var/log` : System log files  
`/var/spool` : Spooling directories for printers, mail, dns  
`/var/tmp` : More temporary space (preserved between reboots)

## Symbolic links

---

- Commonly used to reorganize a subtree, or provide multiple points of access to a file
- “Soft links” -- record path information, but not actual file
- Created by **ln -s**, deleted with **rm**
- Can contain absolute or relative path
  - # **ln -s ./.. parent**
  - # **ln -s /etc/mime.types .mime.types**
- First arg is recorded, not resolved until use

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Character and block device files /dev/

---

- Allow programs to communicate with hardware
  - When kernel gets request that refers to device file, it is handed off to the device driver
- Character (raw) device files: drivers do i/o buffering
- Block device files: handle i/o in large chunks
- Characterized by major (which driver) and minor (which device) device numbers

```
crw-rw---- 1 root lp  6,  0 Jan 30 2003 /dev/lp0
```
- Created with **mknod** and deleted by **rm**
  - Usually managed automatically by system

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## File attributes

---

- Every file has 12 mode bits (four octal values of 3 bits each)
- First three bits:
  - 4000 – setuid
  - 2000 – setgid
  - 1000 – sticky bit
- On a directory, means only the owner of the file, directory, or superuser can delete or rename files
- Keeps /tmp more private and secure

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Sockets & pipes

---

- Local domain sockets
  - Sockets provide connections between processes
  - Local/UNIX domain sockets are only accessible through the filesystem
  - Only used by processes involved in connection
  - Created with socket, deleted by **rm** or **unlink**
  - Used by X Windows, syslog, and printing system
- Named pipes
  - FIFO files that allow communication between processes on same host
  - Created with **mknod** and deleted with **rm**

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison



# Permission bits

---

- Nine permission bits
  - User:owner read, write, execute
    - 400, 200, 100
  - Group read, write, execute
    - 40, 20, 10
  - Other:world read, write execute
    - 4, 2, 1
- Ability to delete or rename is controlled by permissions on directory

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison

## Examples

---

```
-rwxr-xr-x  3 root  root  63555 Mar 13  2002 /bin/gzip
crw-rw-r--  1 root  root  4,   0 Aug  4  2003 /dev/tty0
```

- **chmod** changes permissions
- **chown** changes ownership and group
  - # **chown -R user:group /home/user**
- **umask**
  - Set shell parameters to control default permissions
  - **umask 027** gives everything to owner, forbids writes to group, and gives nothing to other users
  - Usually set in `/etc/profile` or `/etc/csh.login`

Spring 2016

CSE 265: System and Network Administration

©2004-2016 Brian D. Davison